

# Lab 1: Stock Price Prediction Analytics using Snowflake & Airflow

Abhinita Sanabada (018320874)

October 3, 2025

## Abstract

We implemented a secure, reproducible stock analytics workflow using Airflow, `yfinance`, and Snowflake. An ETL DAG ingests OHLCV data into `RAW.STOCK_PRICES`. A second DAG trains a Snowflake-native `SNOWFLAKE.ML.FORECAST` model and writes predictions to `MODEL.FORECASTS`. A final table, `ANALYTICS.FINAL_PRICES_FORECAST`, unions actuals and forecasts for downstream visualization. All Snowflake credentials (account, user, password, role, warehouse, database) are stored only in Airflow Connections, and pipeline parameters are managed via Airflow Variables.

## 1 Problem Statement

Build an end-to-end analytics pipeline that:

1. Extracts daily OHLCV for selected tickers via `yfinance`.
2. Forecasts daily close prices using Snowflake's built-in ML forecasting.
3. Unifies actuals and forecasts in a single analytics table.
4. Uses Airflow for orchestration, SQL/Python transactions for correctness, and Airflow Connections/Variables for secure configuration.
5. Produces reproducible runs, screenshots, and a public code repository.

**Success criteria:** both DAGs succeed; RAW, MODEL, ANALYTICS populated; final table supports plotting Actual vs. Forecast; screenshots and repo links are provided.

## 2 System Architecture

### Overview

We use three schemas inside `USER_DB_CATFISH`: **RAW** (ingest), **MODEL** (predictions), **ANALYTICS** (consumption). Two DAGs orchestrate the flow:

- **DAG #1 `yfinance_etl`:** downloads OHLCV for `stock_symbols`, MERGEs into `RAW.STOCK_PRICES`

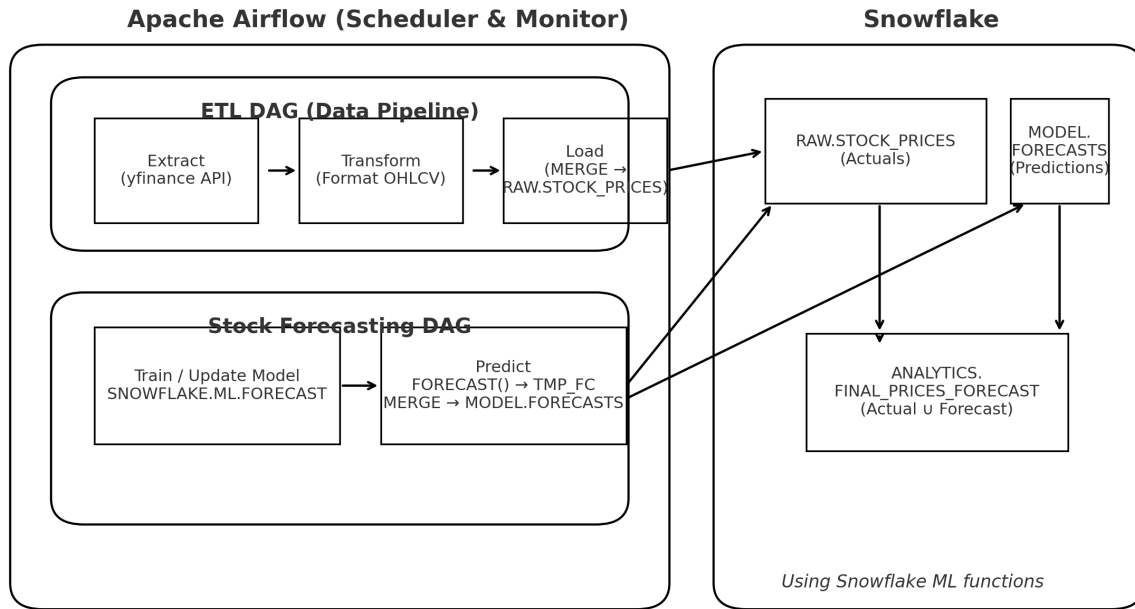


Figure 1: Architecture diagram

- **DAG #2 ml\_forecast**: trains/updates `SNOWFLAKE.ML.FORECAST` on multi-series history; writes predictions to `MODEL.FORECASTS`; *unions* actuals + forecasts into `ANALYTICS.FINAL_PRICES_FORECAST`

## Architecture Diagram

### Screenshots (to be included)

- **Figure 2**: Airflow DAGs list showing both DAGs present.
- **Figure 3**: `yfinance_etl` Grid/Graph view with successful run.
- **Figure 4**: `ml_forecast` Grid/Graph view with successful run.

## 3 Data Model

All objects live in database `USER_DB.CATFISH` (warehouse: `CATFISH_QUERY_WH`, both configured via Airflow Connection). Schemas: `RAW`, `MODEL`, `ANALYTICS`.

### RAW.STOCK\_PRICES

```

SYMBOL STRING NOT NULL; TS TIMESTAMP_NTZ NOT NULL; OPEN FLOAT; HIGH
FLOAT; LOW FLOAT; CLOSE FLOAT; ADJ_CLOSE FLOAT; VOLUME NUMBER(38,0);
LOAD_TS TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP();

```

**Primary key**: (SYMBOL, TS).

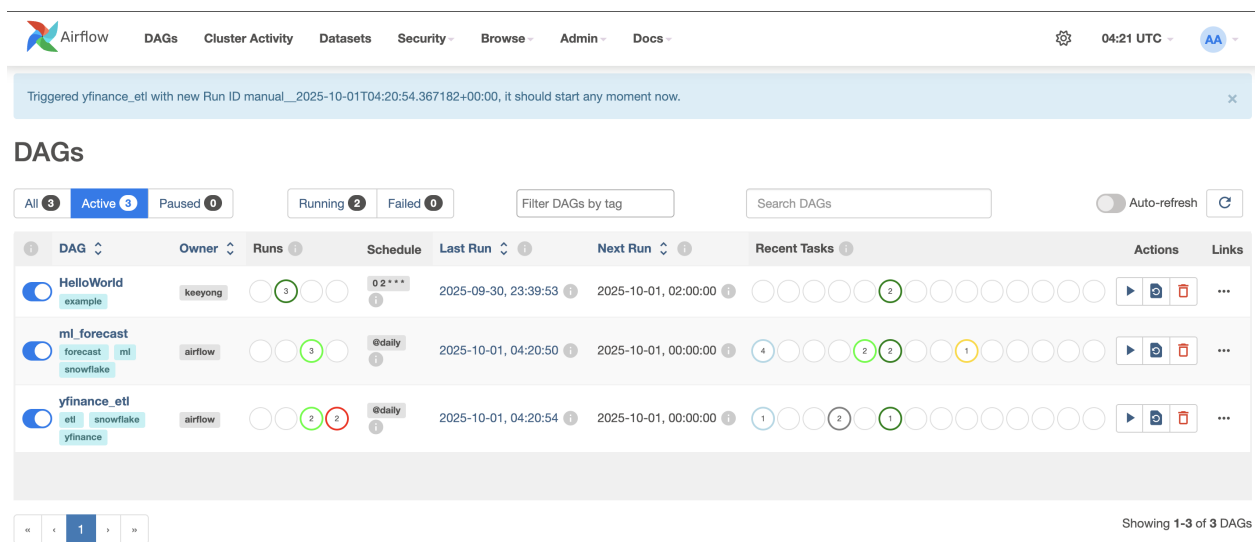


Figure 2: Airflow dags screen shot

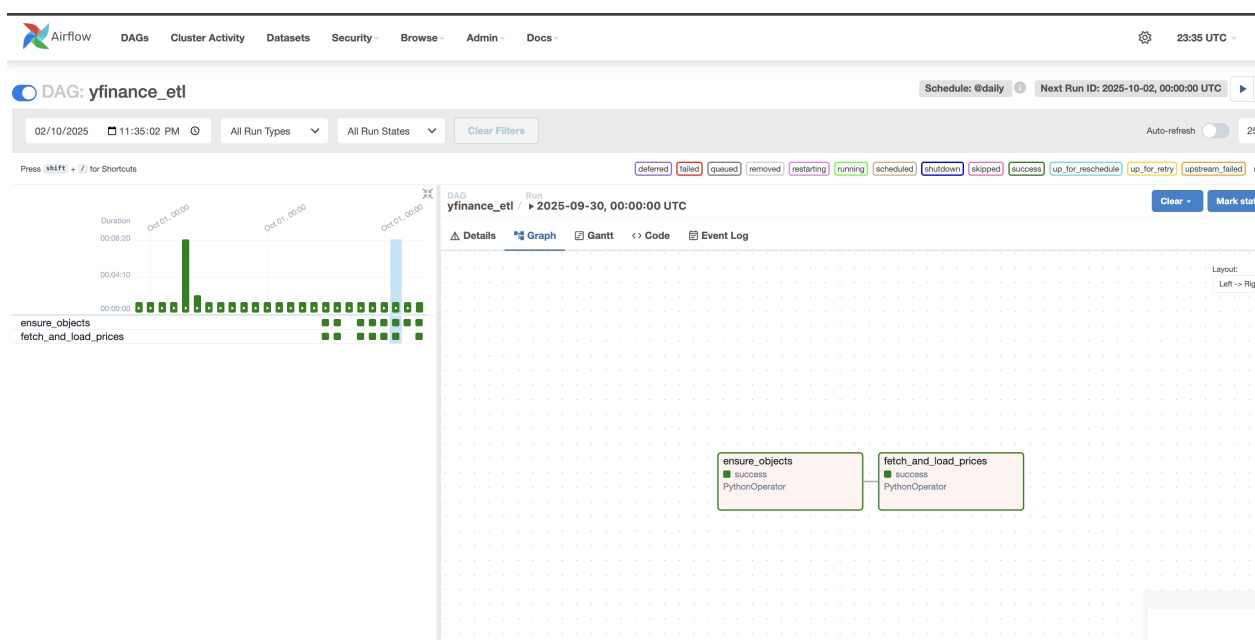


Figure 3: yfinance etl Grid/Graph




Figure 4: ml forecast Grid/Graph

## MODEL.FORECASTS

SYMBOL STRING NOT NULL; TS DATE NOT NULL; PREDICTED\_CLOSE FLOAT NOT NULL; MODEL\_NAME STRING NOT NULL; TRAINED\_AT TIMESTAMP\_NTZ NOT NULL; HORIZON\_D NUMBER(5,0) NOT NULL; LOAD\_TS TIMESTAMP\_NTZ DEFAULT CURRENT\_TIMESTAMP  
**Primary key:** (SYMBOL, TS, MODEL\_NAME).

## ANALYTICS.FINAL\_PRICES\_FORECAST

SYMBOL STRING NOT NULL; TS DATE NOT NULL; CLOSE FLOAT; SOURCE STRING NOT NULL; MODEL\_NAME STRING; LOAD\_TS TIMESTAMP\_NTZ DEFAULT CURRENT\_TIMESTAMP  
**Primary key:** (SYMBOL, TS, SOURCE); SOURCE ∈ {ACTUAL, FORECAST}.

## One-time Bootstrap DDL

Listing 1: Bootstrap DDL (run once).

```

1 BEGIN;
2 CREATE SCHEMA IF NOT EXISTS RAW;
3 CREATE SCHEMA IF NOT EXISTS MODEL;
4 CREATE SCHEMA IF NOT EXISTS ANALYTICS;
5
6 CREATE TABLE IF NOT EXISTS RAW.STOCK_PRICES (
7     SYMBOL STRING NOT NULL, TS TIMESTAMP_NTZ NOT NULL,
8     OPEN FLOAT, HIGH FLOAT, LOW FLOAT, CLOSE FLOAT, ADJ_CLOSE FLOAT,
9     VOLUME NUMBER(38,0),

```

```

9   LOAD_TS TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP(),
10  CONSTRAINT PK_STOCK_PRICES PRIMARY KEY (SYMBOL, TS)
11 );
12
13 CREATE TABLE IF NOT EXISTS MODEL.FORECASTS (
14   SYMBOL STRING NOT NULL, TS DATE NOT NULL, PREDICTED_CLOSE FLOAT
15   NOT NULL,
16   MODEL_NAME STRING NOT NULL, TRAINED_AT TIMESTAMP_NTZ NOT NULL,
17   HORIZON_D NUMBER(5,0) NOT NULL,
18   LOAD_TS TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP(),
19   CONSTRAINT PK_FORECASTS PRIMARY KEY (SYMBOL, TS, MODEL_NAME)
20 );
21
22 CREATE TABLE IF NOT EXISTS ANALYTICS.FINAL_PRICES_FORECAST (
23   SYMBOL STRING NOT NULL, TS DATE NOT NULL, CLOSE FLOAT,
24   SOURCE STRING NOT NULL, MODEL_NAME STRING, LOAD_TS TIMESTAMP_NTZ
25   DEFAULT CURRENT_TIMESTAMP(),
26   CONSTRAINT PK_FINAL PRIMARY KEY (SYMBOL, TS, SOURCE)
27 );
28
29 COMMIT;

```

## 4 Implementation

### 4.1 Airflow Connections & Variables

We created a Snowflake Connection **snowflake\_catfish** with account, user, password, role, default warehouse CATFISH\_QUERY\_WH, and database USER\_DB\_CATFISH. **No secrets in code.** Variables:

- stock\_symbols: JSON list, e.g., ["AAPL", "MSFT", "TSLA"].
- lookback\_days: e.g., 365.
- forecast\_horizon\_days: e.g., 14.
- target\_schema\_raw=RAW, target\_schema\_model=MODEL, target\_schema\_analytics=ANALYTICS

### 4.2 DAG #1: yfinance\_etl (ETL)

Python downloads OHLCV for stock\_symbols over the last lookback\_days and MERGES into RAW.STOCK\_PRICES. Credentials/DB/WH/role come from the Airflow Connection at run-time (via snowflake.connector). We use a transactional pattern (commit/rollback).

### 4.3 Create Snowflake Connection in Airflow UI

**Step 1:** Open the Airflow web UI (e.g., <http://localhost:8080>).

**Step 2:** Navigate to **Admin** → **Connections**.

**Step 3:** Click + (*Add a new record*).

**Step 4:** Fill the form:

- **Conn Id:** snowflake\_catfish
- **Conn Type:** Snowflake
- **Login:** your Snowflake *username*
- **Password:** your Snowflake *password*
- **Extra (JSON):** paste the following and adjust values:

```
{
  "account":    "abcd-xy123",
  "warehouse": "COMPUTE_WH",
  "database":   "YOUR_DB",
  "schema":    "RAW",
  "role":       "SYSADMIN"
}
```

**Step 5:** Click **Test** (top-right). If successful, click **Save**.

#### Notes.

- The DAGs access this via `BaseHook.get_connection("snowflake_catfish")` and read fields from `c.extra_dejson`.
- If your ML pipeline expects a different default schema (e.g., MODEL), either change it here or pass a schema argument in code.

## 4.4 Create Airflow Variables

**Step 1:** Go to **Admin** → **Variables**.

**Step 2:** Click + to add each key/value below (use exact keys):

**Runtime DML (executed by the DAG after staging rows):**

### SQL

Listing 2: ETL MERGE (executed each run).

```
1
2 BEGIN;
3 CREATE TEMP TABLE TMP_LOAD (
4   SYMBOL STRING, TS TIMESTAMP_NTZ, OPEN FLOAT, HIGH FLOAT, LOW FLOAT
5   ,
6   CLOSE FLOAT, ADJ_CLOSE FLOAT, VOLUME NUMBER(38,0)
```

```

6 );
7 -- Python inserts many rows into TMP_LOAD via executemany(...)
8
9 MERGE INTO RAW.STOCK_PRICES AS t
10 USING TMP_LOAD AS s
11   ON  t.SYMBOL = s.SYMBOL
12     AND t.TS      = s.TS
13 WHEN MATCHED THEN UPDATE SET
14   OPEN=s.OPEN, HIGH=s.HIGH, LOW=s.LOW, CLOSE=s.CLOSE,
15   ADJ_CLOSE=s.ADJ_CLOSE, VOLUME=s.VOLUME, LOAD_TS=CURRENT_TIMESTAMP
16   ()
17 WHEN NOT MATCHED THEN INSERT (
18   SYMBOL, TS, OPEN, HIGH, LOW, CLOSE, ADJ_CLOSE, VOLUME
19 ) VALUES (
20   s.SYMBOL, s.TS, s.OPEN, s.HIGH, s.LOW, s.CLOSE, s.ADJ_CLOSE, s.
21   VOLUME
22 );
23 COMMIT;

```

## 4.5 DAG #2: ml\_forecast (Snowflake ML) + Final Union

This DAG uses only `SnowflakeOperator`. It trains/updates a multi-series model via `SNOWFLAKE.ML.FORECAST`, stages horizon-wide predictions, MERGEs them into `MODEL.FORECASTS`, and rebuilds `ANALYTICS.FINAL_RESULTS` by unioning ACTUAL and FORECAST rows.

### Model Training, Forecasting & Upsert

Listing 3: Snowflake ML model + forecast + upsert.

```

1 BEGIN;
2 USE SCHEMA MODEL;
3
4 WITH symbols AS (
5   SELECT value::string AS symbol
6   FROM TABLE(FLATTEN(input => PARSE_JSON('{{ var.value.stock_symbols
7   }}}}'))
8 ),
9 training_data AS (
10  SELECT
11    TO_VARIANT(sp.SYMBOL) AS SERIES,
12    sp.TS,
13    sp.CLOSE
14  FROM RAW.STOCK_PRICES sp
15  JOIN symbols s ON s.symbol = sp.SYMBOL

```

```

15 WHERE sp.TS >= DATEADD('day', -{{ var.value.lookback_days |
    default('365', true) }}, CURRENT_TIMESTAMP())
16 )
17
18 CREATE OR REPLACE SNOWFLAKE.ML.FORECAST PRICE_FORECASTER (
19 INPUT_DATA      => SYSTEM$QUERY_REFERENCE($$ SELECT SERIES, TS,
    CLOSE FROM training_data $$),
20 SERIES_COLNAME   => 'SERIES',
21 TIMESTAMP_COLNAME => 'TS',
22 TARGET_COLNAME   => 'CLOSE',
23 CONFIG_OBJECT    => {{ '{{' }} 'method':'fast','on_error':'skip'
    '{{ '}}' }}
24 );
25
26 CREATE OR REPLACE TEMP TABLE TMP_FC AS
27 SELECT
28     SERIES::STRING                AS SYMBOL,
29     CAST(TS AS DATE)              AS TS,
30     FORECAST                      AS PREDICTED_CLOSE,
31     'SNOWFLAKE_ML'               AS MODEL_NAME,
32     CURRENT_TIMESTAMP()           AS TRAINED_AT,
33     {{ var.value.forecast_horizon_days | default('14', true) }}::
        NUMBER AS HORIZON_D
34 FROM TABLE(PRICE_FORECASTER!FORECAST(
35     FORECASTING_PERIODS => {{ var.value.forecast_horizon_days |
        default('14', true) }}
36 ));
37
38 MERGE INTO MODEL.FORECASTS AS t
39 USING TMP_FC AS s
40 ON t.SYMBOL      = s.SYMBOL
41 AND t.TS         = s.TS
42 AND t.MODEL_NAME = s.MODEL_NAME
43 WHEN MATCHED THEN UPDATE SET
44     PREDICTED_CLOSE = s.PREDICTED_CLOSE,
45     TRAINED_AT      = s.TRAINED_AT,
46     HORIZON_D       = s.HORIZON_D,
47     LOAD_TS         = CURRENT_TIMESTAMP()
48 WHEN NOT MATCHED THEN INSERT (
49     SYMBOL, TS, PREDICTED_CLOSE, MODEL_NAME, TRAINED_AT, HORIZON_D
50 ) VALUES (
51     s.SYMBOL, s.TS, s.PREDICTED_CLOSE, s.MODEL_NAME, s.TRAINED_AT, s.
        HORIZON_D
52 );
53 COMMIT;

```

## Final Union Build (ACTUAL FORECAST)



Listing 4: Rebuild ANALYTICS final table.

```
1 BEGIN;
2 USE SCHEMA ANALYTICS;
3
4 WITH symbols AS (
5     SELECT value::string AS symbol
6     FROM TABLE(FLATTEN(input => PARSE_JSON('{{ var.value.stock_symbols
7         }}}'))
8 )
9 TRUNCATE TABLE ANALYTICS.FINAL_PRICES_FORECAST;
10
11 -- ACTUALS from RAW
12 INSERT INTO ANALYTICS.FINAL_PRICES_FORECAST (SYMBOL, TS, CLOSE,
13     SOURCE, MODEL_NAME)
14 SELECT
15     sp.SYMBOL,
16     CAST(sp.TS AS DATE) AS TS,
17     sp.CLOSE,
18     'ACTUAL' AS SOURCE,
19     NULL AS MODEL_NAME
20 FROM RAW.STOCK_PRICES sp
21 JOIN symbols s ON s.symbol = sp.SYMBOL;
22
23 -- FORECASTS from MODEL
24 INSERT INTO ANALYTICS.FINAL_PRICES_FORECAST (SYMBOL, TS, CLOSE,
25     SOURCE, MODEL_NAME)
26 SELECT
27     f.SYMBOL,
28     f.TS,
29     f.PREDICTED_CLOSE AS CLOSE,
30     'FORECAST' AS SOURCE,
31     f.MODEL_NAME
32 FROM MODEL.FORECASTS f
33 JOIN symbols s ON s.symbol = f.SYMBOL;
34 COMMIT;
```

## 4.6 Transactions & Error Handling

- **DAG #1:** `snowflake.connector` uses `conn.autocommit(False)` and wraps DDL/DML in `try/except` with `commit()` on success and `rollback()` on failure.
- **DAG #2:** Each `SnowflakeOperator` task uses `BEGIN; ...COMMIT;` so the entire step is atomic.

## 5 Results

After triggering `yfinance_etl` and then `ml_forecast`, we validated:

- `RAW.STOCK_PRICES` contains daily OHLCV for each ticker.
- `MODEL.FORECASTS` contains `forecast_horizon_days` predictions per ticker with `MODEL_NAME='SNOWFLAKE_ML'`.
- `ANALYTICS.FINAL_PRICES_FORECAST` holds both ACTUAL and FORECAST rows.

### Validation Queries

```
1 SELECT 'RAW' AS t, COUNT(*) c FROM RAW.STOCK_PRICES
2 UNION ALL SELECT 'MODEL', COUNT(*) FROM MODEL.FORECASTS
3 UNION ALL SELECT 'ANALYTICS', COUNT(*) FROM ANALYTICS.
   FINAL_PRICES_FORECAST;
4
5 SELECT SYMBOL, MIN(TS) AS min_ts, MAX(TS) AS max_ts, COUNT(*) AS n
6 FROM ANALYTICS.FINAL_PRICES_FORECAST
7 GROUP BY SYMBOL
8 ORDER BY SYMBOL, min_ts;
```

## 6 Discussion

**Pros.** Snowflake-side training avoids heavy Python deps and simplifies daily orchestration. Airflow manages schedule, retries, and secret handling via Connections/Variables. Transactions ensure atomic loads and clear failure modes.

### create the Airflow Connection and Variables

```
SYMBOL STRING NOT NULL; TS DATE NOT NULL; CLOSE FLOAT; SOURCE STRING
NOT NULL; MODEL_NAME STRING; LOAD_TS TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP (
```

**Primary key:** (SYMBOL, TS, SOURCE); SOURCE ∈ {ACTUAL, FORECAST}.

Airflow Setup: Snowflake Connection & Variables

#### Required

- `stock_symbols` (JSON):  
["AAPL", "MSFT", "TSLA"]
- `lookback_days`: 365
- `target_schema_raw`: RAW

### Optional (recommended)

- `forecast_horizon_days: 14`
- `target_schema_model: MODEL`
- `target_schema_analytics: ANALYTICS`

### Tips.

- Ensure JSON variables are valid (double quotes, no trailing commas).
- Access in code with `Variable.get("stock_symbols");` parse JSON if needed.
- UI changes to Connections/Variables take effect immediately; new DAG code may need a short scheduler refresh window.

## 7 Quick Checklist

- `snowflake_catfish` exists and **Test** passes.
- `stock_symbols`, `lookback_days`, and `target_schema_raw` are present.
- DAG code references the same *Conn Id* and variable names.

## 8 Common Issues & Fixes

- **“The conn\_id isn’t defined”**: the Connection ID in UI doesn’t match the DAG.
- **Auth failures**: wrong account/role/warehouse in Extra.
- **JSON errors in Variables**: re-save with valid JSON (use the code blocks above).

## 9 Conclusion

The pipeline meets the lab goals: ETL → Snowflake ML Forecast → union table; clean Airflow orchestration with transactions; secure configuration using Connections/Variables.

## References

- yfinance (PyPI): <https://pypi.org/project/yfinance/>
- Apache Airflow: <https://airflow.apache.org/>
- Snowflake Documentation: <https://docs.snowflake.com/>